



Convergence Comparison: Best Methods Across All Benchmarks

Demonstrating Oracle Hierarchy: Hessian > Gradient > No Gradient


Benchmark A: Linear Regression




Benchmark B: Toy Neural Network




Benchmark C: Rosenbrock ($\kappa \approx 2508$)



Benchmark A: Suboptimality Gap



Benchmark B: Suboptimality Gap



Benchmark C: Suboptimality Gap

