

Adagrad vs SGD on a Sparse Gradient Problem (coord 0 updates 10× rarer)

