



Q1: Gradient Norm vs. Iteration
(Benchmark C, Rosenbrock)



Effective Step Size: Adagrad vs. RMSprop
(Benchmark C)



Adagrad decays to 0
(frozen learning)

RMSprop stabilises
(adaptive)