


Q2: Mini-Batch SGD vs Full-Batch GD (Benchmark A)

Per-Epoch Convergence
(SGD with more updates dominates)



Per-Gradient-Evaluation Convergence
(full GD is more efficient per unit cost)

