



Newton's Method vs. Gradient Descent: Suboptimality $f(x_k) - f^*$
Newton's quadratic convergence vs GD's linear rate

Benchmark A
(Quadratic, $\kappa \approx 1$)



Benchmark B
(NN loss, $\kappa \approx 7$)



Benchmark C
(Rosenbrock, $\kappa \approx 2504$)

