


Variance Reduction: SGD vs SVRG vs SAGA (Benchmark A)

SGD vs SVRG vs SAGA on Benchmark A
(Linear Regression, $m = 1000$, batch=40)



SGD Gradient Variance vs SVRG Convergence
(Variance $\rightarrow 0$ explains linear rate)

