

Mini-Batch SGD: Gradient Variance vs Batch Size
(Benchmark A, at $\hat{\theta}_{\text{OLS}}$, $n = 200$ trials)

