

Mini-batch gradient variance vs. batch size
(Benchmark A, $\theta_0 = [0, 0]$)

